

Hierarchical classification of breast cancer molecular subtypes from discrete aCGH copy-number profiles

Alexandros Michailidis (2903034)¹, Antonie Wagner (2903383)², Christos Botos (2878553)¹, and Yan Qiao (2874296)¹

¹Bioinformatics and Systems Biology Master’s Programme, Vrije Universiteit Amsterdam, The Netherlands

²Computer Science Master’s Programme, Vrije Universiteit Amsterdam, The Netherlands

Abstract

Motivation: Breast cancer molecular subtypes (HER2+, HR+, Triple Negative) require distinct therapeutic strategies, yet classification from discrete array CGH copy-number profiles ($n = 100$, 2,834 regions) poses a high-dimensional, low-sample-size challenge. Wessels et al. (2005) showed that simple classifiers with univariate filtering match complex pipelines on continuous expression data, but whether this holds for discrete aCGH data with multi-class structure has not been tested.

Results: A 2×2 factorial experiment crossing two feature selectors (Kruskal-Wallis, Elastic Net) with two classifiers (Nearest Mean, Random Forest) revealed that Random Forest outperformed Nearest Mean in flat 3-class classification (BA = 0.791 vs 0.631). However, HER2+ recall was consistently high across all flat pipelines (0.764-0.998), suggesting that HR+ vs Triple Negative confusion was the true performance bottleneck. A hierarchical decomposition (Stage 1: HER2+ vs rest; Stage 2: HR+ vs TN) reversed the ranking, with Nearest Mean outperforming Random Forest in Stage 2. The best pipeline, hierarchical EN+NMC with plateau ensembling, achieved a combined 3-class balanced accuracy of 0.840 and an expected validation accuracy of 48 of 57 samples.

Availability: Code and predictions are provided in the submission archive.

1 Introduction

Breast cancer comprises molecular subtypes, including HER2-positive (HER2+), hormone receptor-positive (HR+), and triple-negative (TN), that differ in prognosis and require distinct therapeutic strategies (Perou et al., 2000; Sørlie et al., 2001). Array Comparative Genomic Hybridization (aCGH) enables genome-

wide detection of DNA copy-number (CN) alterations (Pinkel and Albertson, 2005); after segmentation and calling (Venkatraman and Olshen, 2007; van de Wiel et al., 2007, 2011), it produces high-dimensional discrete profiles suitable for classification. The present study classifies breast cancer subtypes from such CN profiles using 100 labelled training samples and 57 unlabelled validation samples, a typical high-dimensional, low-sample-size challenge.

Wessels et al. (2005) demonstrated that simple classifiers with univariate feature filtering match or outperform complex pipelines on continuous expression data (van ’t Veer et al., 2002), and Haury et al. (2011) confirmed that univariate selection achieves competitive accuracy while offering better feature stability. However, these findings pertain to continuous data; whether they extend to discrete aCGH profiles and multi-class settings has not been tested. We therefore ask: does the Wessels principle hold for discrete CN data, and how do feature redundancy and multi-class structure affect classifier selection? To answer this, we employ a 2×2 factorial design comparing simple and complex feature selection and classification methods via nested cross-validation (CV), and explore a hierarchical decomposition motivated by the observation that HER2+ samples are trivially separable while HR+ vs TN confusion is the true bottleneck.

2 Methods

2.1 Data description

The training set comprises $n = 100$ samples profiled on a 244K-probe aCGH platform, summarized into 2,834 genomic regions with discrete CN calls: loss (-1), normal (0), gain (+1), or amplification (+2). Classes are HER2+ ($n = 32$), HR+ ($n = 36$), and TN ($n = 32$). A validation set of 57 unlabelled samples

is provided. Coordinates use hg18/NCBI Build 36. The analysis was implemented in Python 3.10 with `scikit-learn` 1.7.2 (Pedregosa et al., 2011).

2.2 Label-free region merging

Adjacent genomic regions showed strong spatial autocorrelation (median Pearson $r = 0.92$ within 0.1 Mb). Following the label-free CGHregions approach of van de Wiel and van Wieringen (2007), we merged adjacent regions with $r > 0.8$ by taking the median CN per sample, reducing the feature space from 2,834 to 273 segments (90.4% reduction). The threshold was chosen because 91% of adjacent pairs exceeded $r = 0.8$, and nested CV tests confirmed that omitting merging degraded performance by 7 pp ($p = 0.008$, Wilcoxon). Merging is entirely label-free and was performed before CV splits; the merging map was saved and applied to validation data at prediction time to ensure consistent segment boundaries. This is distinct from the label-dependent leakage described by Ambroise and McLachlan (2002).

2.3 Experimental design

Following Wessels et al. (2005), we adopted a 2×2 factorial design crossing two feature selectors with two classifiers (Figure 1).

Feature selection. The Kruskal-Wallis (KW) test (Kruskal and Wallis, 1952) ranks features by their H -statistic across the three subtypes; the top $k \in \{5, 10, 15, 20, 30, 50, 75, 100\}$ are retained. As a multivariate alternative, Elastic Net (EN) multinomial logistic regression (Zou and Hastie, 2005) evaluates features jointly; the top $k \in \{5, 10, 20, 50\}$ (by summed absolute coefficients) are retained, with $C \in \{10^{-3}, \dots, 10^2\}$ and `l1_ratio` $\in \{0.1, 0.3, 0.5, 0.7, 0.9\}$ tuned in the inner CV loop.

Classifiers. The Nearest Mean Classifier (NMC) assigns samples to the closest class centroid with optional shrinkage $\delta \in \{\text{None}, 0.1, 0.2, 0.5\}$ (Wessels et al., 2005). Random Forest (RF) uses 500 trees with balanced class weights (Breiman, 2001), tuning `max_features` $\in \{\text{sqrt}, 0.3\}$ and `min_samples_leaf` $\in \{1, 5\}$. Crossing the two selectors with two classifiers yields four pipelines with 32 (KW) or 800 (EN) inner CV configurations per fold.

2.4 Hierarchical classifier

HER2+ recall was near-perfect across RF pipelines (0.994-0.998; Table 1), while HR+ vs TN confusion was the bottleneck across all flat pipelines. A binary KW ranking on HR+ and TN samples surfaced chr12q and chr5q features (Figure 2B) instead of the chr17

ERBB2 signal dominating 3-class ranking, motivating a two-stage architecture.

Stage 1 classifies HER2+ vs rest using a fixed KW+RF model with $k = 5$ chr17q12 features at the ERBB2 amplicon ($H = 73.9$, $p = 9 \times 10^{-17}$), achieving BA = 1.000 in all 1,000 outer folds. Stage 2 classifies the remaining samples as HR+ or TN using one of the four pipeline variants with full inner CV tuning on non-HER2+ samples only, accessing the chr12q/chr5q signal that 3-class ranking obscures (Figure 1).

2.5 Cross-validation scheme

We employed nested, stratified, repeated 5-fold CV (Wessels et al., 2005; Ambroise and McLachlan, 2002). The outer loop estimates generalisation performance (50 repeats $\times$ 5 folds = 250 evaluations for flat; $200 \times 5 = 1,000$ for hierarchical). The inner 5-fold loop tunes hyperparameters, with all feature selection performed strictly within inner folds to prevent leakage. Performance is measured by balanced accuracy (Brodersen et al., 2010). The workflow is shown in Figure 1.

2.6 Plateau ensembling

Inner CV folds in Stage 2 contain only ~ 11 samples, insufficient to reliably rank 800 EN configurations: 303 distinct top-1 selections appeared across 1,000 folds, with negative inner/outer-test correlation ($r = -0.378$). Plateau ensembling addresses this by retaining all configurations within one pooled SD of the best mean score (up to 15) and averaging their class probabilities for prediction. The threshold was fixed before observing outer-fold scores. Leave-one-repeat-out analysis confirmed plateau stability (fewer than 1 of 15 members changed per omitted repeat). The plateau gain may be slightly inflated by indirect inner/outer CV dependence, though this analysis suggests the effect is small.

2.7 Statistical comparison and final model

Pipeline differences were assessed using the Nadeau-Bengio corrected t -test (Nadeau and Bengio, 2003), which accounts for non-independence in repeated CV, and bootstrap CIs (10,000 resamples) (Bouckaert and Frank, 2004). The final submitted model (EN+NMC with plateau ensembling) was retrained on all 100 samples with fixed plateau membership and the same region-merging map applied to validation data.




Figure 1: Nested CV workflow. The outer loop (200 repeats $\times$ 5 folds) estimates performance; the inner loop tunes hyperparameters. Stage 1 separates HER2+ (BA1 = 1.000); Stage 2 classifies HR+ vs TN. Feature selection occurs inside each inner fold.

3 Results

3.1 Flat 2x2 experiment

In the flat setting, RF outperformed NMC regardless of feature selector: KW+RF achieved BA = 0.791 vs 0.631 for KW+NMC (Table 1). HER2+ recall was high across all pipelines (0.764-0.998), but TN recall ranged from only 0.507 to 0.641 (Figure 3B, left), identifying HR+ vs TN confusion as the bottleneck.

3.2 Hierarchical classifier

The hierarchical decomposition reversed the classifier ranking: NMC outperformed RF in Stage 2 (HR+ vs TN), with bootstrap CIs excluding zero for all NMC-vs-RF pairs. EN+NMC(P) achieved the best

BA = 0.840 and BA2 = 0.759 (Table 1; Figure 3). The Nadeau-Bengio test found no significant pairwise differences, consistent with its conservative nature for repeated CV.

3.3 Plateau ensembling

Plateau ensembling improved EN+NMC from BA2 = 0.729 to 0.759 (+3.0 pp; Table 1). The gain scaled with grid size: EN (800 configs) gained +3.0 pp vs KW (32 configs) at +1.4 pp, explaining why EN+NMC(P) overtook KW+NMC(P) despite a weaker base BA2 (0.729 vs 0.736).

Table 1: Pipeline performance summary. BA is 3-class balanced accuracy; BA2 is Stage 2 balanced accuracy (HR+ vs TN). BA1 = 1.0 for all hierarchical pipelines. Values are mean $\pm$ SD of per-repeat means. Per-class columns show mean $\pm$ SD of per-repeat recalls. Med. k is the median number of features selected across outer folds. Best mean values per column in bold.

Experiment	Pipeline	BA	BA2	HER2+	HR+	TN	Med. k
Flat	KW+NMC	0.631 $\pm$ 0.031	–	0.764 $\pm$ 0.069	0.623 $\pm$ 0.064	0.507 $\pm$ 0.079	15
	KW+RF	0.791 $\pm$ 0.025	–	0.994 $\pm$ 0.014	0.736 $\pm$ 0.056	0.641 $\pm$ 0.056	75
	EN+NMC	0.746 $\pm$ 0.032	–	0.909 $\pm$ 0.035	0.698 $\pm$ 0.047	0.630 $\pm$ 0.066	10
	EN+RF	0.770 $\pm$ 0.029	–	0.998 $\pm$ 0.008	0.696 $\pm$ 0.064	0.616 $\pm$ 0.058	20
Hier. (base)	KW+NMC	0.824 $\pm$ 0.031	0.736 $\pm$ 0.046	1.000 $\pm$ 0.000	0.760 $\pm$ 0.056	0.712 $\pm$ 0.068	30
	KW+RF	0.814 $\pm$ 0.029	0.721 $\pm$ 0.044	1.000 $\pm$ 0.000	0.748 $\pm$ 0.058	0.692 $\pm$ 0.057	30
	EN+NMC	0.819 $\pm$ 0.026	0.729 $\pm$ 0.040	1.000 $\pm$ 0.000	0.729 $\pm$ 0.051	0.730 $\pm$ 0.058	20
	EN+RF	0.794 $\pm$ 0.031	0.691 $\pm$ 0.046	1.000 $\pm$ 0.000	0.726 $\pm$ 0.058	0.655 $\pm$ 0.070	50
Hier. (plateau)	KW+NMC(P)	0.834 $\pm$ 0.026	0.750 $\pm$ 0.040	1.000 $\pm$ 0.000	0.773 $\pm$ 0.047	0.727 $\pm$ 0.064	100
	EN+NMC(P)	0.840 $\pm$ 0.022	0.759 $\pm$ 0.032	1.000 $\pm$ 0.000	0.777 $\pm$ 0.045	0.741 $\pm$ 0.047	81
	Ensemble	0.838 $\pm$ 0.025	0.757 $\pm$ 0.037	1.000 $\pm$ 0.000	0.778 $\pm$ 0.046	0.735 $\pm$ 0.056	–

3.4 Single best biomarker region

The chr17q12 ERBB2 amplicon (KW $H = 73.9$, $p = 9 \times 10^{-17}$) enabled perfect Stage 1 separation using five features, consistent with its established role as a clinical biomarker (Slamon et al., 1987). For Stage 2, discriminative signal resided on chr12q and chr5q (Figure 2).

3.5 Accuracy estimate

The expected number of correct predictions was estimated as $\sum_c \hat{r}_c \cdot (n_c/N) \cdot 57$, where $\hat{r}_c$ is the mean per-class recall across 200 repeats and n_c/N the training class proportion. With HER2+ = 1.000, HR+ = 0.777, TN = 0.741, this yielded $47.7 \approx 48$. Sensitivity analysis under alternative class proportions

produced estimates of 46-48; 48 was submitted.

3.6 Improvements after draft

Three changes improved accuracy after the draft. (1) The flat KW+RF (BA = 0.791) was replaced by a hierarchical decomposition, improving the best BA to 0.824 (+3.3 pp) and reversing the classifier ranking: NMC outperformed RF in Stage 2, consistent with the Wessels principle on a well-structured subproblem. (2) Plateau ensembling raised EN+NMC from BA = 0.819 to 0.840 (+2.1 pp) and BA2 from 0.729 to 0.759 (+3.0 pp). (3) The estimate was updated from BA $\times$ 57 to a per-class weighted formula, increasing from 45 to 48.




Figure 2: Feature selection stability for EN+NMC(P) Stage 2. (A) Selection frequency across 1,000 outer folds for all 273 segments. (B) Top 15 regions, color-coded by chromosome. Chr12 and chr5 dominate, distinct from the chr17 ERBB2 signal in Stage 1.




Figure 3: Model comparison. **(A)** BA distributions per pipeline (50 repeats flat, 200 hierarchical). All hierarchical variants outperform the best flat pipeline. **(B)** Confusion matrices for KW+RF (flat, left) and EN+NMC(P) (hierarchical, right), pooled over all outer folds.

4 Discussion

The Wessels principle is conditional on problem structure. In the flat 3-class setting, RF outperformed NMC (BA = 0.791 vs 0.631; Figure 3A), apparently contradicting findings on continuous expression data. After hierarchical decomposition removed the trivially separable HER2+ class, NMC outperformed RF in the binary HR+ vs TN problem. The HER2+ class acted as a confounder: its perfect separability inflated RF’s advantage in the 3-class setting. In both settings, feature selector choice mattered less than classifier choice, consistent with Wessels et al. (2005).

Limitations include the small sample size ($n = 100$, 68 for Stage 2) and discrete CN representation. Approximately 10 samples were consistently misclassified (>80% across 1,000 test appearances; Figure 3B), sug-

gesting the BA2 ≈ 0.759 ceiling reflects irreducible noise rather than classifier capacity. Future work could integrate expression data for multi-omic classification and test whether the hierarchical NMC advantage generalises to larger cohorts. In conclusion, simple classifiers can outperform complex ones on discrete aCGH data, but only when multi-class structure is decomposed to match the underlying biology.

Author contributions

C. Botos: EDA, region merging, hierarchical nested CV, plateau ensembling, cpu-cluster. **A. Wagner:** Flat nested CV, KW+NMC, KW+RF, research question conceptualization. **A. Michailidis:** Flat nested CV, EN+NMC, EN+RF, standalone EN. **Y. Qiao:** Final training, estimate computation. All authors contributed to intellectual

design and report writing.

References

- C. Ambroise and G. J. McLachlan. Selection bias in gene extraction on the basis of microarray gene-expression data. *Proceedings of the National Academy of Sciences*, 99(10):6562–6566, 2002. doi: 10.1073/pnas.102102699.
- R. R. Bouckaert and E. Frank. Evaluating the replicability of significance tests for comparing learning algorithms. In *Advances in Knowledge Discovery and Data Mining*, pages 3–12, 2004. doi: 10.1007/978-3-540-24775-3_3.
- L. Breiman. Random forests. *Machine Learning*, 45(1): 5–32, 2001. doi: 10.1023/A:1010933404324.
- K. H. Brodersen, C. S. Ong, K. E. Stephan, and J. M. Buhmann. The balanced accuracy and its posterior distribution. In *Proceedings of the 20th International Conference on Pattern Recognition*, pages 3121–3124, 2010. doi: 10.1109/ICPR.2010.764.
- A.-C. Haury, P. Gestraud, and J.-P. Vert. The influence of feature selection methods on accuracy, stability and interpretability of molecular signatures. *PLoS ONE*, 6(12):e28210, 2011. doi: 10.1371/journal.pone.0028210.
- W. H. Kruskal and W. A. Wallis. Use of ranks in one-criterion variance analysis. *Journal of the American Statistical Association*, 47(260):583–621, 1952. doi: 10.1080/01621459.1952.10483441.
- C. Nadeau and Y. Bengio. Inference for the generalization error. *Machine Learning*, 52(3):239–281, 2003. doi: 10.1023/A:1024068626366.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- C. M. Perou, T. Sørlie, M. B. Eisen, M. van de Rijn, S. S. Jeffrey, C. A. Rees, J. R. Pollack, D. T. Ross, H. Johnsen, L. A. Akslen, O. Fluge, A. Pergamenschikov, C. Williams, S. X. Zhu, P. E. Lønning, A.-L. Børresen-Dale, P. O. Brown, and D. Botstein. Molecular portraits of human breast tumours. *Nature*, 406(6797):747–752, 2000. doi: 10.1038/35021093.
- D. Pinkel and D. G. Albertson. Array comparative genomic hybridization and its applications in cancer. *Nature Genetics*, 37:S11–S17, 2005. doi: 10.1038/ng1569.
- D. J. Slamon, G. M. Clark, S. G. Wong, W. J. Levin, A. Ullrich, and W. L. McGuire. Human breast cancer: correlation of relapse and survival with amplification of the HER-2/neu oncogene. *Science*, 235(4785):177–182, 1987. doi: 10.1126/science.3798106.
- T. Sørlie, C. M. Perou, R. Tibshirani, T. Aas, S. Geisler, H. Johnsen, T. Hastie, M. B. Eisen, M. van de Rijn, S. S. Jeffrey, T. Thorsen, H. Quist, J. C. Matese, P. O. Brown, D. Botstein, P. E. Lønning, and A.-L. Børresen-Dale. Gene expression patterns of breast carcinomas distinguish tumor subclasses with clinical implications. *Proceedings of the National Academy of Sciences*, 98(19): 10869–10874, 2001. doi: 10.1073/pnas.191367098.
- M. A. van de Wiel and W. N. van Wieringen. CGHregions: dimension reduction for array CGH data with minimal information loss. *Cancer Informatics*, 3:55–63, 2007. doi: 10.1177/117693510700300031.
- M. A. van de Wiel, K. I. Kim, S. J. Vosse, W. N. van Wieringen, S. M. Wilting, and B. Ylstra. CGHcall: calling aberrations for array CGH tumor profiles. *Bioinformatics*, 23(7):892–894, 2007. doi: 10.1093/bioinformatics/btm030.
- M. A. van de Wiel, R. Brosens, B. Ylstra, and W. N. van Wieringen. Preprocessing and downstream analysis of microarray DNA copy number profiles. *Briefings in Bioinformatics*, 12(1):10–21, 2011. doi: 10.1093/bib/bbq004.
- L. J. van ’t Veer, H. Dai, M. J. van de Vijver, Y. D. He, A. A. M. Hart, M. Mao, H. L. Peterse, K. van der Kooy, M. J. Marton, A. T. Witteveen, G. J. Schreiber, R. M. Kerkhoven, C. Roberts, P. S. Linsley, R. Bernards, and S. H. Friend. Gene expression profiling predicts clinical outcome of breast cancer. *Nature*, 415(6871):530–536, 2002. doi: 10.1038/415530a.
- E. S. Venkatraman and A. B. Olshen. A faster circular binary segmentation algorithm for the analysis of array CGH data. *Bioinformatics*, 23(6):657–663, 2007. doi: 10.1093/bioinformatics/btl646.
- L. F. A. Wessels, M. J. T. Reinders, A. A. M. Hart, C. J. Veenman, H. Dai, Y. D. He, and L. J. van’t Veer. A protocol for building and evaluating predictors of disease state based on microarray data. *Bioinformatics*, 21(19): 3755–3762, 2005. doi: 10.1093/bioinformatics/bti429.
- H. Zou and T. Hastie. Regularization and variable selection via the elastic net. *Journal of the Royal Statistical Society: Series B*, 67(2):301–320, 2005. doi: 10.1111/j.1467-9868.2005.00503.x.